

# Draft: Building an elaborator using extensible constraints

Bohdan Liesnikov, Jesper Cockx

## ABSTRACT

Dependently-typed languages are used for statically enforcing properties of programs and for enabling type-driven development. While there’s been a lot of work to find the right theoretical foundations for the core type theory, the implementations have been studied to a smaller extent. In particular, theoretical works rarely consider the plethora of features that exist in bigger languages like Agda, so the developers have to make a lot of adaptations in the implementations, making them rather ad-hoc, in particular in the elaborator. And when these features appear organically over time they can lead to codebases that are hard to maintain. We present a new design for the elaboration of dependently-typed languages based on the idea of an open datatype for constraints and plugin system for solvers to tackle these issues. This allows for a more compact base elaborator implementation while enabling extensions to the type system.

## ACM Reference Format:

Bohdan Liesnikov, Jesper Cockx. 2023. Draft: Building an elaborator using extensible constraints. In *Proceedings of ACM Conference (Conference’17)*. ACM, New York, NY, USA, 11 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

## 1 INTRODUCTION

Strongly-typed languages allow us to catch certain classes of bugs at compile-time by checking the implementation against the type-signature. When the types are provided by the user it can be viewed as a form of specification, constraining the behaviour of the programs. This comes with the benefit of more static guarantees but with an increased toll on the user to supply more precise information about the program. Since type is part of the specification we can use this information to infer parts of our program, following an idea that Connor McBride aptly worded as “Write more types and fewer programs.” (PTOOP 2022; McBride 2005, chap. 2.1) Some examples here include overloaded functions in Java, implicits in Scala, and type classes in Haskell.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

*Conference’17, July 2017, Washington, DC, USA*  
© 2023 Association for Computing Machinery.  
ACM ISBN 978-x-xxxx-xxxx-x/YY/MM. . . \$15.00  
<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

In dependently-typed languages like Agda (Norell 2007; The Agda Team 2022), Coq (The Coq Development Team 2022) or Idris (Brady 2013) the types can be much more precise. This gives us an opportunity to infer even larger parts of our programs from the type in the process of elaboration. Techniques here include implicit arguments in Agda, implicit coercions in Coq, and tactic arguments in Idris. The inference or “solving” can be not only automatic but also interactive or partially automatic. For example, holes in Agda, proof obligations and canonical structures (Mahboubi and Tassi 2013) in Coq, and holes in Haskell (Koppel et al. 2022). All of these mechanisms use different solvers and have various degrees of extensibility. They are usually not isolated from each other and can therefore interact in the way the user expects them to (for example, in this case between implicits and instances (Agda users 2018)).

In all of these examples, the solvers evolved organically over time together with the language. Coq historically it’s been struggling with similar issues in the elaboration: for example, Canonical Structures didn’t get to be properly documented for 15 years (Mahboubi and Tassi 2013). Agda experimented with features baked into the core of the type system, like sized types which brought their own solver infrastructure (Abel and Winterhalter 2016). Lean 4 is aiming to allow the users to develop new surface-level features (Leonardo de Moura and Sebastian Ullrich 2021) using elaboration monads (Leonardo de Moura and Ullrich 2021), but Lean 3 was built in a more conventional way (Leonardo de Moura et al. 2015). These solvers operate on constraints generated by the type-checker, where the constraints say something about as-of-yet unknown parts of the program, which are referred to as metavariables, also known as “existential variables” (The Coq Development Team 2022, chap. 2.2.1).

While the features above may be specific to the language, a common need is to check two terms for equality, dubbed “conversion checking”, which goes hand in hand with unification. This code is also heavily used throughout the compiler, for inference of implicit arguments and for general type-checking, making it sensitive towards changes and hard to maintain and debug. It is hard enough to understand that breaking changes are often discovered only when ran against a large existing project on CI, like `cubical` or `stdlib` for Agda or `unimath` for Coq. Unification problems can be rendered as constraints and that’s a view we will use throughout the paper.

We propose a new architecture for an extensible elaborator for dependently-typed languages. The idea is to provide an API that allows users to tap into the elaboration procedure with their custom solvers that can manipulate metavariables and

constraints placed on them. We aim to further the understanding the art and science of implementing dependently-typed languages, making the implementations of different features of the language more independent. This design separates the what the solvers are doing from the when, making it explicit what are the interaction points between them where the developer has to pay attention.

Contributions:

- We propose a new design blueprint for a language that is extensible with new constraints and their solvers. It supports type classes, implicit arguments, implicit coercions, context arguments (as implemented in Scala) and tactic arguments.
- We provide a suite of solvers in lieu of common solvers like the conversion checker in Agda.
- We suggest a new view on metavariables as communication channels for the solvers.
- We have a work-in-progress implementation of a prototype of a dependently-typed language with an extendable unifier, implicit arguments, type classes,

## 2 UNIFICATION, CONSTRAINT-BASED ELABORATION AND DESIGN CHALLENGES

Constraints have been an integral part of compilers for strongly-typed languages for a while (Odersky, Sulzmann, and Wehr 1999). For example, both Haskell (Vytiniotis et al. 2011) and Agda (Norell 2007, chap. 3) use constraints extensively. In the former case, they are even reflected and can be manipulated by the user (Orchard and Schrijvers 2010; GHC development team 2022, chap. 6.10.3). This has proved to be a good design decision for GHC, as is reflected, for example in a talk by Peyton Jones (2019), as well as in a few published sources (Vytiniotis et al. 2011; Peyton Jones et al. 2007).

In the land of dependently-typed languages constraints are much less principled. Agda has a family of constraints<sup>1</sup> that grew organically, currently counting 17 constructors. Idris technically has constraints<sup>2</sup> with the only two constructors being equality constraint for two terms and for two sequences of terms. The same<sup>3</sup> holds for Lean. These languages either use constraints in a very restricted, single use-case manner – for unification, namely – or in an unprincipled way. In our view, more methodical approach to constraints will result in a more robust elaborator as a whole.

<sup>1</sup>We shorten the links in footnotes to paths in the repository, the source code can be found at [github.com/agda/agda/blob/v2.6.2.2/./src/full/Agda/TypeChecking/Monad/Base.hs#L1064-L1092](https://github.com/agda/agda/blob/v2.6.2.2/./src/full/Agda/TypeChecking/Monad/Base.hs#L1064-L1092)

<sup>2</sup>[./src/Core/UnifyState.idr](https://github.com/agda/agda/blob/v2.6.2.2/./src/Core/UnifyState.idr) at [github.com/idris-lang/Idris2/blob/e673d0](https://github.com/idris-lang/Idris2/blob/e673d0)

<sup>3</sup>[./src/Lean/Match/Basic.lean#L161](https://github.com/leanprover/lean4/blob/0a031f3./src/Lean/Match/Basic.lean#L161) at [github.com/leanprover/lean4/blob/0a031f3](https://github.com/leanprover/lean4/blob/0a031f3)

In this section we go over three typical challenges that come up while building a compiler for a dependently-typed language and the way they are usually solved. We cover unification of the base language and the complexity of managing the state of the unifier in section 2.1. Then we take a look at different kinds of implicit arguments and their implementation in section 2.2. We briefly touch on the problem of extending the unifier in section 2.3. Finally, we talk about the ideas behind the new design in section 2.4.

### 2.1 Conversion checking in the presence of a meta-variables

As mentioned in the introduction, in the process of type-checking a program we need to compare terms and it is the job of the unifier, which is notoriously hard to implement. The complexity stems from the desire of compiler writers to implement the most powerful unifier, while being limited by the fact that it higher-order unification is undecidable in general. Some of this complexity is unavoidable, but we can manage it better by splitting it up into small modular components. In practice, this means that one doesn't have to fit together an always-growing conversion checker but can instead write different cases separately. We rely on the constraint solver machinery to distribute the problems to the fitting solvers.

An example from Agda's conversion checker is `compareAs` function<sup>4</sup> which provides type-driven conversion checking. The function is almost 90 lines long, and yet the vast majority of it is special cases of metavariables. This function calls the `compareTerm'` function<sup>5</sup>, which itself is 130 lines. `compareTerm'` calls the `compareAtom` function<sup>6</sup>. Which itself is almost 200 lines of code. Each of the above functions implements part of the “business logic” of the conversion checker. But each of them contains a lot of code dealing with bookkeeping related to metavariables and constraints:

1. They have to throw and catch exceptions, driving the control flow of the unification.
2. They have to compute blocking tags that determine when a postponed constraint is retried.
3. They have to deal with cases where either or both of the sides equation or its type are either metavariables or terms whose evaluation is blocked on some metavariables.

This code is also heavily used throughout the type-checker: either as direct functions `leqType` when type-checking terms, `compareType` when type-checking applications, or as raised constraints `ValueCmp` and `SortCmp` from `equalTerm` while checking applications or definitions, `ValueCmpOnFace` from `equalTermOnFace` again while checking applications. And at

<sup>4</sup>[./src/full/Agda/TypeChecking/Conversion.hs#L146-L218](https://github.com/agda/agda/blob/v2.6.2.2/./src/full/Agda/TypeChecking/Conversion.hs#L146-L218)

<sup>5</sup>[./src/full/Agda/TypeChecking/Conversion.hs#L255-L386](https://github.com/agda/agda/blob/v2.6.2.2/./src/full/Agda/TypeChecking/Conversion.hs#L255-L386)

<sup>6</sup>[./src/full/Agda/TypeChecking/Conversion.hs#L419-L675](https://github.com/agda/agda/blob/v2.6.2.2/./src/full/Agda/TypeChecking/Conversion.hs#L419-L675)

the same it is unintuitive and full of intricacies as indicated by multiple comments<sup>7</sup>.

We would like the compiler-writer to separate the concerns of managing constraints, blockers from the actual logic of the comparison function. In fact, if we zoom in on the `compareAtom` function, the core can be expressed in about 20 lines<sup>8</sup> of simplified code, stripping out size checks, cumulativity, polarity, and forcing. Which is precisely what we'd like the developer to write.

```
case (m, n) of
  (Lit l1, Lit l2) | l1 == l2 -> return ()
  (Var i es, Var i' es') | i == i' -> do
    a <- typeOfBV i
    compareElims [] [] a (var i) es es'
  (Con x ci xArgs, Con y _ yArgs) | x == y -> do
    t' <- conType x t
    compareElims t' (Con x ci []) xArgs yArgs
  ...
```

The functions described above are specific to Agda but in other major languages we can find similar problems with unifiers being large modules that are hard to understand. The sizes of modules with unifiers are as follows: Idris (1.5kloc), Lean (1.8kloc), Coq (1.8kloc). For Haskell, which is not a dependently-typed language yet but does have a constraints system (Peyton Jones 2019), this number is at 2kloc.

## 2.2 Type-checking function application in the presence of implicit arguments

During the type-checking of function application there may be different kinds of arguments to process, for example, instance arguments, implicit arguments, or tactic arguments. If we start from a simple case of type-checking an application of a function symbol to regular arguments, every next extension requires to be handled in a special case.

Take Agda – when checking an application during the insertion of implicit arguments<sup>9</sup> we already have to carry the information on how the argument will be resolved and then create a specific kind of metavariables<sup>10</sup> (Norell 2007, chap. 3) – for each of those cases.

Idris 2 (The Idris Team 2021, chap. 13.1) in the case of `auto` variables has to essentially do inline the search procedure through a chain of elaboration function calls (`checkApp` to `checkAppWith` to `checkAppWith'`) to `makeAutoImplicit`<sup>11</sup>. This can accommodate interfaces (or type classes), but one can imagine that if a different kind of implicit was added, like tactic arguments, or Canonical Structures we'd have to

<sup>7</sup>./src/full/Agda/TypeChecking/Conversion.hs#L430-L431,

./src/full/Agda/TypeChecking/Conversion.hs#L521-L529

<sup>8</sup>./src/full/Agda/TypeChecking/Conversion.hs#L530-L579

<sup>9</sup>./src/full/Agda/TypeChecking/Implicit.hs#L99-L127

<sup>10</sup>./src/full/Agda/TypeChecking/Implicit.hs#L131-L150

<sup>11</sup>./src/TTImp/Elab/App.idr#L224-L241

inline the search again, requiring a non-trivial modification to the elaboration mechanism.

While the codebases above show that it is certainly possible to extend languages with new features if the language wasn't written with extensibility in mind this can lead to rather ad-hoc solutions. Instead is of handling every kind of metavariable in a separate way we'd like to uniformly dispatch a search for the solution, which is then handled by the constraint solvers. We can achieve this by creating metavariables for the unknown terms and then raising a constraint for the meta containing the type of the meta. Then this constraint can be matched on by the appropriate solver based on the type and the elaborator for the application of a function doesn't have to know anything about the implicits at all. The only thing we require is that the elaboration of the argument is called with the type information available. This corresponds to how in bidirectional typing function application is done in the inference mode but the arguments are processed in checking mode.

```
inferType (App t1 t2) = do
  (et1, Pi tyA tyB) <- inferType t1
  et2 <- checkType t2 tyA
  return (App et1 et2, subst tyB et2)
```

```
checkType (Implicit) ty = do
  m <- createMetaTerm
  raiseConstraint $ FillInTheMeta m ty
  return m
```

This metavariable in its own turn gets instantiated by a fitting solver. The solvers match the shape of the type that metavariable stands for and handle it in a case-specific manner: instance-search for type classes, tactic execution for a tactic argument.

If it is a regular implicit, however, the only solver that's needed is a trivial one that checks that the metavariable has been instantiated indeed. This is because a regular implicit should be instantiated by a unification problem encountered at some point later. This serves as a guarantee that all implicits have been filled in.

Let us go through an example of the elaboration process for a simple term:

```
plus : {A : Type} -> {{PlusOperation A}}
      -> (a : A) -> (b : A) -> A
```

```
instance PlusNat : PlusOperation Nat where
  plus = plusNat
```

```
two = plus 1 1
```

We will step through the elaboration of the term `two`.

1. First, the pre-processor eliminates the implicits and type-class arguments. We end with the following declarations:

291  
292  
293  
294  
295  
296  
297  
298  
299  
300  
301  
302  
303  
304  
305  
306  
307  
308  
309  
310  
311  
312  
313  
314  
315  
316  
317  
318  
319  
320  
321  
322  
323  
324  
325  
326  
327  
328  
329  
330  
331  
332  
333  
334  
335  
336  
337  
338  
339  
340  
341  
342  
343  
344  
345  
346  
347  
348

```

plus : (impA : Implicit Type)
      -> TypeClass PlusOperation (deImp impA)
      -> (a : deImp impA) -> (b : deImp impA)
      -> deImp impA

```

```

PlusNat = Instance {
  class = PlusOperation Nat,
  body = {plus = plusNat}}

```

```
two = plus _ _ 1 1
```

2. We go into the elaboration of `two` now. The elaborator applies `inferType (App t1 t2)` rule four times and `checkType (Implicit) ty` twice on the two placeholders. The output of the elaborator is

```
two = plus ?_1 ?_2 1 1
```

And the state of the elaborator contains four more constraints:

```

C1: FillInTheTerm ?_1 (Implicit Type)
C2: FillInTheTerm ?_2 (TypeClass
    PlusOperation (deImp ?_1))
C3: EqualityConstraint ?_1 Nat Type
C4: EqualityConstraint ?_1 Nat Type

```

The first two correspond to implicit arguments. The latter two are unification problems rendered into constraints.

3. Now we step into the constraint-solving world. First, the unifier solves the latter two, instantiating `?_1` to `Nat`. Next, the typeclass resolution launches a search for the instance, resolving `?_2` to the `PlusNat` instance. Finally, `C1` is discarded as solved since `?_1` is already instantiated to `Nat`.

## 2.3 Extending unification

While writing a unifier is hard enough as it is, at times the developers might want to give their users ability to extend the unification procedure.

Canonical Structures (Saibi 1999; Mahboubi and Tassi 2013) was already mentioned as it is in the overlap between type classes and unification hints. Adding it to a language that doesn't support them requires an extension of unification algorithm with a rule that says that projection from a canonical structure is an injective function (Mahboubi and Tassi 2013, equation 1).

One could also provide means to do so manually in a more general case, by allowing users to declare certain symbols as injective. This is one of the features requested by the Agda users (Agda users 2023).

Another example of this can be adding rules of associativity and commutativity to the unifier, as described in the thesis by Holten (2023). It required 2000 lines of code added while

relying on rewrite rules<sup>12</sup> We would like to make changes such as this more feasible.

## 2.4 What is our design bringing into the picture?

The examples above show that when building a dependently-typed language while the core might be perfectly elegant and simple, the features that appear on top of it complicate the design.

One can also observe that while the code above might rely on constraints, the design at large doesn't put at the centre of the picture and instead is primarily seen as a gadget. To give a concrete example, Agda's constraint solver<sup>13</sup> relies on the type-checker to call it at the point where it is needed and has to be carefully engineered to work with the rest of the codebase. To give a concrete example, functions `noConstraints` or `dontAssignMetas` rely on specific behaviour of the constraint solver system and are used throughout the type-checker. `abortIfBlocked`, `reduce` and `catchConstraint/patternViolation` force the programmer to make a choice between letting the constraint system handle blockers or doing it manually. These things are known to be brittle and pose an increased mental overhead when making changes.

Our idea for a new design is to shift focus more towards the constraints themselves:

1. We give a stable API for raising constraints that can be called by the type-checker, essentially creating an "ask" to be fulfilled by the solvers. This is not dissimilar to the idea of mapping object-language unification variables to host-language ones as done by Guidi, Coen, and Tassi (2017), view of the "asks" as a general effect (Bauer, Haselwarter, and Petković 2020, ch. 4.4) or communication between actors (Allais et al. 2022).
2. Make the language more modular we make constraints an extensible data type in the style of Swierstra (2008) and give an API to define new solvers with the ability to specify what kinds of constraints they can solve.

In the examples in this paper, we follow the bidirectional style of type-checking, but in practice, the design decisions are agnostic of the underlying system, as long as it adheres to the principle of stating the requirements on terms in terms of raising a constraint and not by, say, pattern-matching on a concrete term representation.

From a birds-eye view the architecture looks as depicted in Figure 1

<sup>12</sup>[github.com/LHolten/agda-commassoc/tree/definitional-commutativity](https://github.com/LHolten/agda-commassoc/tree/definitional-commutativity)

<sup>13</sup>[src/full/Agda/TypeChecking/Constraints.hs#L251-L301](https://src/full/Agda/TypeChecking/Constraints.hs#L251-L301)



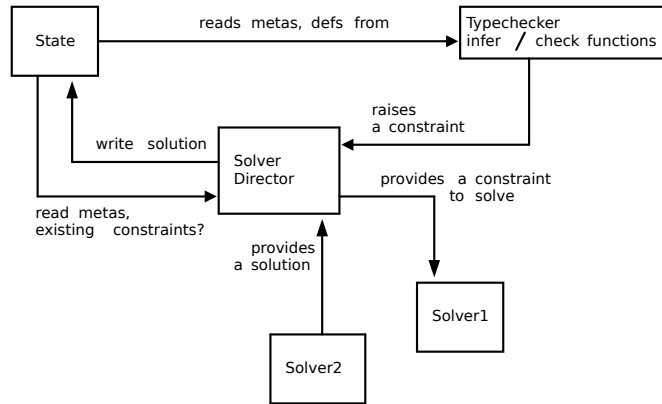


Figure 1: Architecture diagram

In the diagram type-checker is precisely the part that implements syntax-driven traversal of the term. It can raise a constraint that gets registered by the Solver Director. Solver Director then is exactly the component that dispatches solvers on the appropriate constraints and constitutes our main contribution. All of the components have some read access to the state, including Solver which might e.g. verify that there are no additional constraints on the meta.

### 3 DEPENDENTLY-TYPED CALCULUS AND BIDIRECTIONAL TYPING

In this section, we describe the core of the type system we implement. We take pi-forall (Weirich 2022) as a basis for the system and add metavariables to it. However, for all other purposes, we leave the core rules intact and therefore, the core calculus too.

#### 3.1 Basic language and rules

This is dependently-typed calculus that includes Pi, Sigma and indexed inductive types.

Here's an almost complete surface-language term data type:

```
data Term =
  -- type of types Type
  Type
  -- variables x
  | Var TName
  -- abstraction \x. a
  | Lam (Bind TName Term)
  -- application a b
  | App Term Arg
  -- function type (x : A) -> B
  | Pi Epsilon Type (Bind TName Type)
  -- Sigma-type { x : A | B }
  | Sigma Term (Bind TName Term)
  | Prod Term Term
```

```
| LetPair Term (Bind (TName, TName) Term)
-- Equality type a = b
| TyEq Term Term
| Refl
| Subst Term Term
| Contra Term
-- inductive datatypes
| TCon TName [Arg] -- types (fully applied)
| DCon DName [Arg] -- terms (fully applied)
| Case Term [Match]
-- metavariables
| MetaVar MetaClosure
```

Equality isn't defined as a regular inductive type, but is instead built-in with the user getting access to the type and term constructor, but not able to pattern-matching on it, instead getting a `subst` primitive of type  $(A \ x) \rightarrow (x=y) \rightarrow A \ y$  and `contra` of type  $\text{forall } A. \text{ True} = \text{False} \rightarrow A$ .

On top of the above we include indexed inductive datatypes and case-constructs for their elimination. Indexed inductive datatypes are encoded using a well-known trick as parameterised inductive datatypes with an equality argument constraining the index.

#### 3.2 Syntax traversal

We implement the core of the elaborator as a bidirectional syntax traversal, raising a constraint every time we need to assert something about the type.

This includes the expected use of unification constraints, like in case we enter check-mode while the term should be inferred:

```
checkType tm ty = do
  (etm, ty') <- inferType tm
  constrainEquality ty' ty I.Type
  return $ etm
```

Any time we want to decompose the type provided in checking mode:

```
checkType (S.Lam lam) ty = do
  mtyA <- createMetaTerm
  mtX <- createUnknownVar
  mtyB <- extendCtx (I.TypeSig (I.Sig mtX mtyA))
    (createMetaTerm)
  let mbnd = bind mtX mtyB
  let metaPi = I.Pi mtyA mbnd

  constrainEquality ty metaPi I.Type
  -- rest of the traversal can now use mtyA and mbnd
  ...
```

At certain points we have to raise a constraint which has an associated continuation. Like for checking the type of a data constructor – the part of the program that comes as an

argument to `CA.constraintConAndFreeze` will be suspended (or “blocked”) until the constraint has been resolved.

```
checkType t@(S.DCon c args) ty = do
  elabpromise <- createMetaTerm
  CA.constraintConAndFreeze ty $ do
    mty <- SA.substMetas ty
    case mty of
      (I.TCon tname params) -> do
        t
        ...
      _ -> ...
```

## 4 CONSTRAINTS AND UNIFICATION

The datatype of constraints is open which means the user can write a plugin to extend it. However, we provide a few out of the box to be able to type-check the base language.

For the purposes of the base language it suffices to have the following.

- Two terms of certain type must be equal, encodes unification problems  
*-- two terms given should be equal*  
`data EqualityConstraint e =`  
`EqualityConstraint Syntax.Term Syntax.Term`  
`Syntax.Type Syntax.MetaVarId`

In the process of unification we also need to do an occurs-check (Abel and Pientka 2011, sec. 3.1), encoded as follows:

```
data OccursCheck e =
  OccursCheck Syntax.Term Syntax.Type
  [Syntax.TName]
```

- Ensures that a metavariable is resolved eventually:  
*-- this terms has to be filled in*  
`data FillInTheTerm e =`  
`FillInTheTerm Syntax.Term Syntax.Type`
- Lastly, we provide a constraint which ensures that a term is a type constructor. This could’ve been encoded as a unification problem, but since we don’t have to limit ourselves in the constructors of the constraint datatype and there are no real downsides to factoring a problem out we include it separately:  
*-- the term passed to the constraint*  
*-- should be a type constructor*  
`data TypeConstructorConstraint e =`  
`TypeConstructorConstraint Syntax.Type`

The type-checker raises them supplying the information necessary, but agnostic of how they’ll be solved.

### 4.1 Introduction to the solvers

On the solver side we provide a suite of unification solvers that handle different cases of the problem:

```
-- solves syntactically equal terms
syntacticSolverHandler :: (EqualityConstraint :<: c)
```

```
=> Constraint c -> MonadElab Bool
syntacticSolver :: (EqualityConstraint :<: c)
=> Constraint c -> MonadElab Bool
syntactic :: Plugin
syntactic = Plugin { solver = syntacticSolver
  , handler = syntacticSolverHandler
  ...
}
```

We first define the class of constraints that will be handled by the solver via providing a “handler” – function that decides whether a given solver has to fire. In this case, this amounts to checking that the constraint given is indeed an `EqualityConstraint` and that the two terms given to it are syntactically equal. Then we define the solver itself. Which in this case doesn’t have to do anything except mark the constraint as solved, since we assume it only fires once it’s been cleared to do so by the handler. The reason for this separation between a decision procedure and execution of it is to ensure separation between effectful and costly solving and cheap decision-making that should require only read-access to the state. Finally, we register the solver by declaring it using a plugin interface. This plugin symbol will be picked up by the linker and registered at the runtime.

Similarly, we can define solvers that only work on problems where only one of the sides is a metavariable, `leftMetaSolver` and `rightMetaSolver` of the same type as the syntactic solver above and corresponding handlers and plugins.

Here the job of the solver is not as trivial – it has to check that the type of the other side indeed matches the needed one and then register the instantiation of the metavariable in the state. If both of those steps are successful we can return `True` and the constraint will be marked as solved.

In the cases above we don’t have to worry about the order since the problems they match on don’t overlap. In the case they don’t we can provide priority preferences:

```
complexSolver1 :: Constraint c -> MonadElab Bool
complexHandler1 :: Constraint c -> MonadElab Bool
complexSymbol1 = "complexSolver1"
complex1 = Plugin { ...
  , symbol = complexSymbol1
  , pre = [ unifySolverLS
    , unifySolverRS]
  , suc = []
}
```

```
complexSolver2 :: Constraint c -> MonadElab Bool
complexHandler2 :: Constraint c -> MonadElab Bool
complexSymbol2 = "complexSolver2"
complex2 = Plugin { ...
  , symbol = complexSymbol2
  , pre = [complexSymbol1]
  , suc = []
}
```

At the time of running the compiler, these preferences are loaded into a big pre-order relation for all the plugins, which is then linearised and used to guide the solving procedure.

## 4.2 Unification

We implement a system that is very close to the system implemented by Abel and Pientka (2011) with the exception that every function call in the simplification procedure is now a raised constraint.

For example, the “decomposition of functions” (Abel and Pientka 2011, fig. 2) rule is translated to the following implementation:

```
piEqInjectivitySolver :: (EqualityConstraint <: cs)
=> SolverType cs
piEqInjectivitySolver constr = do
  let (Just (EqualityConstraint (I.Pi a1 b1)
                                (I.Pi a2 b2) _ m)) =
      match @EqualityConstraint constr
  ma <- constrainEquality a1 a2 I.Type
  (x, tyB1, _, tyB2) <- unbind2Plus b1 b2
  let mat = I.identityClosure ma
  mb <- extendCtx (I.TypeSig (I.Sig x e1 mat)) $
      constrainEquality tyB1 tyB2 I.Type
  let mbt = bind x $ I.identityClosure mb
  solveMeta m (I.Pi mat mbt)
  return True
```

The seemingly spurious metavariable `m` serves as an anti-unification (Pfenning 1991) communication channel. When an equality constraint is created we return a metavariable that stands for the unified term. This metavariable is used for unification problems that are created in the extended contexts – in this case second argument of the `Pi`-type, but also when solving equalities concerning two data constructors. We do this to solve the “spine problem” (V́ctor Ĺpez Juan 2021, sec. 1.4) – we operate according to the “well-typed modulo constraints” principle essentially providing a placeholder that is guaranteed to preserve well-typedness in the extended context. In case one of the constraints created in the extended context when solving it we might need to know the exact shape of `ma`. In which case we can block on the metavariable, freezing the rest of the problem until it is instantiated.

As for the actual unification steps, we implement them in a similar fashion to simplification procedure. A big difference is that we have to implement an occurs-check when instantiating metavariables. This also happens through the mechanism of constraints, but this time using the `OccursCheck` constraint mentioned in section 4.

Take a look at the following example, when only the left hand side of the constraint is an unsolved meta:

```
leftMetaSolver :: (EqualityConstraint <: cs)
=> SolverType cs
```

```
leftMetaSolver constr = do
  let (Just (EqualityConstraint t1 t2 _ m)) =
      match @EqualityConstraint constr
  cid <- occursCheck t1 t2
  blockOn (Blockers.All [cid]) $ do
    let (MetaVar (MetaVarClosure m1 c1)) = t1
    (Just ic1) = closure2Subst <$> invertClosure c1
    st2 = subs ts ic1 t2
    solveMeta m1 st2
    solveMeta m st2
  return True
```

Once the occurs-check constraint has been created we block on it and write a solution to the metavariable only when it is actually correct. In case occurs-check fails the constraint simply won’t get solved and therefore the elaboration procedure fails too.

By splitting up the rules into individual, simple solvers we can compartmentalise the complexity of the unifier, making sure that each rule is as decoupled from the others as possible. This doesn’t influence the properties of the system, but doesn’t help to guarantee them either. If one wishes to prove correctness of the unification the same work has to be done as in the usual setting.

## 4.3 Extending unification

Let’s create a simple plugin that makes certain symbols declared by the user injective.

The actual implementation is relatively simple and isn’t dissimilar to the `Pi`-injectivity solver we showed above.

```
userInjectivitySolver :: (EqualityConstraint <: cs)
=> SolverType cs
userInjectivitySolver constr = do
  let (Just EqualityConstraint (I.App (I.Var f) a)
                                (I.App (I.Var g) b) _ m) =
      match @EqualityConstraint constr
  if f == g
  then do
    ifM (queryInjectiveDeclarations f)
    (do
      solveMeta m <=< constrainEquality a b I.Type
      return True)
    (return False)
  else return False
```

where `queryInjectiveDeclarations` simply scans the available declarations for a marker that `f` has been declared injective.

The only big thing left is to make sure that this solver fires at the right time. This can only conflict with the “decomposition of neutrals” rule, so we indicate to the Solver Direction that our plugin should run before the it:

```

813 userInjectivityPlugin :: (EqualityConstraint <:: cs)
814                      => Plugin cs
815 userInjectivityPlugin =
816   Plugin { ...
817     , solver = userInjectivitySolver
818     , symbol = "userInjectivity"
819     , pre = [ unifyNeutralsDecomposition
820              , unificationStartMarkerSymbol]
821     , suc = [unificationStartMarkerSymbol]
822   }

```

This modification doesn't alter the core of the language. However, implementing something like commutativity and associativity unifiers can require modifications in the core, since we two terms that are equal during elaboration have to equal during core type-checking too.

## 5 OPEN DATATYPE OF CONSTRAINTS AND CASE STUDIES

Let's now take a look at what the openness of the constraint datatype buys us in this design. In this section we'll first briefly describe how implicit arguments are implemented and then showcase how to build an extension implementing type-classes into the language.

### 5.1 Implicit arguments

Plan:

- show what we need the pre-processor to produce
  - one option here is what's described in section 1, to translate
 

```

847 def f : {a : A} -> B a -> C
848
849 def f : (A : Implicit Type)
850       -> (a : Implicit (deImp A))
851       -> (b : B (deImp a)) -> C

```
  - the other option is to make `Implicit A` compute to `A` and make solvers not eager to reduce. In that case `deImp` is unnecessary.
- implement tactic arguments with a custom solver for each tactic, like one that just exhaustively searches constructors

### 5.2 Type classes

Plan:

- show what we have to desugar instance and type-class declarations to similar to section 1, but more general
- show the implementation of the solvers
- showcase two alternative definitions of the type-class resolution?
- can we extend this to canonical structures here?

## 6 LIMITATIONS

While our design offers a lot of flexibility, some of the things are challenging to implement. In this section we describe a few examples of extensions that don't quite fit in this framework.

### 6.1 Handling of meta-variables outside of definition sites

After we elaborate a definition there can still be unsolved metas in it. This presents us with a design choice. The first option is to freeze the metas and instantiate them per-use site, essentially allowing for an expression to have multiple types. The second option is to leave them up to be solved later, which might make elaboration less predictable since now the use sites can influence whether a particular definition type-checks. The third option is to report them as an error and exit immediately.

In particular, the second option allows us to incorporate more involved inference algorithms into the system. For example, if we were to implement an erasure inference algorithm as described by Tejiščák (2020), we would have to create metavariable annotations (described as “evars” in the paper) that can be instantiated beyond the definition site.

### 6.2 The language is only as extensible, as the syntax traversal is

Extensibility via constraints allows for a flexible user-specified control flow as soon as we step into the constraints world. But control flow of the main body of the basic language elaborator is fixed by the basic language developer. For example, consider the following simplified lambda-function typechecking function<sup>14</sup> from Agda:

```

checkLambda' cmp b xps typ body target = do
  TelV tel btyp <- telViewUpTo numbinds target
  if size tel < numbinds || numbinds /= 1
  then dontUseTargetType
  else useTargetType tel btyp

```

Here Agda steps away from the bidirectional discipline and infers a (lambda) function if the target type isn't fully known. If in our design the developer chooses to go only with a pure bi-directional style of type-checking inferred lambda functions would be impossible to emulate. That is unless one essentially renders macros and writes their own typechecking case for an inferrable lambda.

In order to gain this extra bit of flexibility we provide `inferType` case for lambdas, even though our base language doesn't use it.

<sup>14</sup>./src/full/Agda/TypeChecking/Rules/Term.hs#L460-L578

actually  
write  
this  
case



### 6.3 Eager reduction and reliance on the pre-processor

This work crucially relies on a pre-processor of some kind, be it macro expansion or some other way to extend the parser with custom desugaring rules. In particular, in order to implement n-ary implicit arguments correctly and easily we need the pre-processor to expand them to the right arity. For coercions, we need to substitute every term  $t$  in the coercible position for `coerce _ t`. This can impact performance.

Alternatively, one can imagine a system where constraint solvers are latching onto non-reduced types and terms in constraints. In that case, we can get around with a trick borrowed from Coq, where `coerce f t` computes to `f t`, but since we typechecked an unreduced application the search will still be launched on the right form.

This also means that constraints can/have to match on unreduced types in the e.g. `FillInTheTherm`

## 7 RELATED WORK

We are certainly not the first ones to try to tackle the extensibility of a language implementation. This section is structured according to the part of the compiler pipeline that allows for extensibility. Dependently-typed language implementations usually consist of at least four parts: parser, elaborator, core type-checker, and proper backend. The back-end part is currently irrelevant to our interests, since for a language to be specified usually means for specification of the core, anything that happens beyond after core doesn't extend the language, but rather tries to preserve its semantics in some form. Therefore we're left with three parts: parser, elaborator and core type-checker.

We see parser or syntax extensibility as a necessary part of an extensible language. This problem has been studied extensively in the past and has a multitude of existing solutions. Macros are one of them and are utilised heavily in various forms in almost all established languages (The Coq Development Team 2022; The Agda Team 2022; Ullrich and de Moura 2020) and can be powerful enough to build a whole language around (Chang et al. 2019).

Core extensibility, on the other hand, appears to be a problem with too many degrees of freedom. Andromeda (Bauer et al. 2018; Bauer, Haselwarter, and Petković 2020) made an attempt at extensible definitional equality but is quite far from a usable dependently-typed language. Agda's philosophy allows developers to experiment with the core but also results in a larger amount of unexpected behaviours. In general, modification of core rules will result in fundamental changes in the type theory, which can break plenty of important properties like soundness or subject reduction.

This leaves us with the question of what can be achieved with the extensibility of an elaborator and why one would need it. This includes type inference, implicit arguments inference,

type classes, tactics, and SMT integration. Elaborators are often structured similarly to the core type-checker, i.e. following a bidirectional discipline of some sort. One can see that in Agda (Norell 2007), Matita (Tassi et al. 2012), or in a paper by Ferreira and Pientka (2014).

Coq (The Coq Development Team 2022) being one of the most popular proof assistants invested a lot of effort into user-facing features: work on tactics like a new tactic engine (Spiwack 2011) and tactic languages (Ltac2 (Pédrot 2019), SSReflect (Gonthier, Mahboubi, and Tassi 2008), etc.), the introduction of a virtual machine for performance (Grégoire and Leroy 2002) and others. However, the implementation is quite hard to extend. One either has to modify the source code, which is mostly limited to the core development team, as seen from the `graph`. Or one has to use Coq plugins system, which is rather challenging, and in the end the complexity of it gave rise to TemplateCoq (Malecha 2014).

Agda introduced a lot of experimental features, but isn't very modular (Robert Estelle 2019), which hinders further change.

Lean introduced elaborator extensions (Leonardo de Moura and Sebastian Ullrich 2021; Ullrich and de Moura 2020). They allow the user to overload the commands, but if one defines a particular elaborator it becomes hard to interleave with others. In a way, this is an imperative view on extensibility.

Idris (Brady 2013; Christiansen and Brady 2016) appeared as a programming language first and proof-assistant second and doesn't provide either a plugin or hook system at all, except for reflection. Idris also focuses on tactics as the main mechanism for elaboration. Metavariables and constraints-wise in contrast with only one kind of meta and constraints in Idris (Brady 2013, chap. 4.2, chap. 4.3.1).

As for the unification,  $\text{Tog}^+$  (V́ctor Ĺpez Juan 2020, 2021) focuses on extending it to the case where two sides of equality might not be of the same type. The main argument against the usage of anti-unification in Agda provided there is that it is bug-prone. On closer inspection, we think that that in Agda, which served as an example of the system where said bugs occur, the problems were stemming from the fact that anti-unification had to be implemented separately from unification, therefore duplicating the code-base. In which case it is indeed hard to keep the two in sync. However, in our case there is no such duplication since unification and anti-unification always move synchronously. Regarding twin types as implemented in  $\text{Tog}^+$  – we don't see a reason why it couldn't be done in this setting too, but since it required a rewrite of the system we decided against it in the interest of time.

## 8 FUTURE WORK

There are some things we leave for future work.

- Implement erasure inference (Tejišćák 2020)?

- Implement Canonical Structures (Mahboubi and Tassi 2013)?
- Rendering of macros as constraints? Map a macro to an implicit term with the right kind of annotation in the type, to get the right expander as an elaboration procedure?
- Mapping constraint solving onto a concurrent execution model. Use LVars here (Kuper 2015) here, similar to what TypOS (Allais et al. 2022) is doing?

## 9 REFERENCES

- Abel, Andreas, and Brigitte Pientka. 2011. "Higher-Order Dynamic Pattern Unification for Dependent Types and Records." In *Typed Lambda Calculi and Applications*, edited by Luke Ong, 10–26. Lecture Notes in Computer Science. Berlin, Heidelberg: Springer. [https://doi.org/10.1007/978-3-642-21691-6\\_5](https://doi.org/10.1007/978-3-642-21691-6_5).
- Abel, Andreas, and Théo Winterhalter. 2016. "An Extension of Martin-Löf Type Theory with Sized Types." In, 2. Novi Sad, Serbia.
- Agda users. 2018. "Performance Regression · Issue #3435 · Agda/Agda." GitHub. December 6, 2018. <https://github.com/agda/agda/issues/3435>.
- . 2023. "Injective for Unification' Pragma · Issue #6546 · Agda/Agda." GitHub. March 23, 2023. <https://github.com/agda/agda/issues/6546>.
- Allais, Guillaume, Malin Altenmuller, Conor McBride, Georgi Nakov, Nordvall Forsberg, and Craig Roy. 2022. "TypOS: An Operating System for Typechecking Actors." In, 3. Nantes, France. [https://types22.inria.fr/files/2022/06/TYPES\\_2022\\_paper\\_31.pdf](https://types22.inria.fr/files/2022/06/TYPES_2022_paper_31.pdf).
- Bauer, Andrej, Gaëtan Gilbert, Philipp G. Haselwarter, Matija Pretnar, and Christopher A. Stone. 2018. "Design and Implementation of the Andromeda Proof Assistant." Schloss Dagstuhl - Leibniz-Zentrum fuer Informatik GmbH, Wadern/Saarbruecken, Germany. <https://doi.org/10.4230/LIPICS.TYPES.2016.5>.
- Bauer, Andrej, Philipp G. Haselwarter, and Anja Petković. 2020. "Equality Checking for General Type Theories in Andromeda 2." In *Mathematical Software – ICMS 2020*, edited by Anna Maria Bigatti, Jacques Carette, James H. Davenport, Michael Joswig, and Timo de Wolff, 253–59. Lecture Notes in Computer Science. Cham: Springer International Publishing. [https://doi.org/10.1007/978-3-030-52200-1\\_25](https://doi.org/10.1007/978-3-030-52200-1_25).
- Brady, Edwin. 2013. "Idris, a General-Purpose Dependently Typed Programming Language: Design and Implementation." *Journal of Functional Programming* 23 (5): 552–93. <https://doi.org/10.1017/S095679681300018X>.
- Chang, Stephen, Michael Ballantyne, Milo Turner, and William J. Bowman. 2019. "Dependent Type Systems as Macros." *Proceedings of the ACM on Programming Languages* 4 (December): 3:1–29. <https://doi.org/10.1145/3371071>.
- Christiansen, David, and Edwin Brady. 2016. "Elaborator Reflection: Extending Idris in Idris." In *Proceedings of the 21st ACM SIGPLAN International Conference on Functional Programming*, 284–97. ICFP 2016. New York, NY, USA: Association for Computing Machinery. <https://doi.org/10.1145/2951913.2951932>.
- Ferreira, Francisco, and Brigitte Pientka. 2014. "Bidirectional Elaboration of Dependently Typed Programs." In *Proceedings of the 16th International Symposium on Principles and Practice of Declarative Programming*, 161–74. PPDP '14. New York, NY, USA: Association for Computing Machinery. <https://doi.org/10.1145/2643135.2643153>.
- GHC development team. 2022. "GHC 9.4.2 User's Guide." GHC manual. August 22, 2022. [https://downloads.haskell.org/ghc/9.4.2/docs/users\\_guide/index.html](https://downloads.haskell.org/ghc/9.4.2/docs/users_guide/index.html).
- Gonthier, Georges, Assia Mahboubi, and Enrico Tassi. 2008. "A Small Scale Reflection Extension for the Coq System." Report. <https://hal.inria.fr/inria-00258384>.
- Grégoire, Benjamin, and Xavier Leroy. 2002. "A Compiled Implementation of Strong Reduction." In *Proceedings of the Seventh ACM SIGPLAN International Conference on Functional Programming*, 235–46. ICFP '02. New York, NY, USA: Association for Computing Machinery. <https://doi.org/10.1145/581478.581501>.
- Guidi, Ferruccio, Claudio Sacerdoti Coen, and Enrico Tassi. 2017. "Implementing Type Theory in Higher Order Constraint Logic Programming." <https://hal.inria.fr/hal-01410567>.
- Holten, Lucas. 2023. "Dependent Type-Checking Modulo Associativity and Commutativity." Master's thesis, Delft, the Netherlands: Delft University of Technology. <https://repository.tudelft.nl/islandora/object/uuid%3A3Af6c2de8-5437-4e47-b37f-6358c04eda9e>.
- Koppel, James, Zheng Guo, Edsko de Vries, Armando Solar-Lezama, and Nadia Polikarpova. 2022. "Searching Entangled Program Spaces." *Proceedings of the ACM on Programming Languages* 6 (August): 91:23–51. <https://doi.org/10.1145/3547622>.
- Kuper, Lindsey. 2015. "Lattice-Based Data Structures for Deterministic Parallel and Distributed Programming." *ProQuest Dissertations and Theses*. PhD thesis, United States – Indiana: Indiana University. <https://www.proquest.com/docview/1723378413/abstract/C592471EAAE94DBFPQ/1>.
- Leonardo de Moura, and Sebastian Ullrich. 2021. "Lean 4 - Metaprogramming." Presented at the Lean together 2021, January 6. <https://leanprover-community.github.io/lt2021/slides/leo-LT2021-meta.pdf>.
- Mahboubi, Assia, and Enrico Tassi. 2013. "Canonical Structures for the Working Coq User." In *Proceedings of the 4th International Conference on Interactive Theorem Proving*, 19–34. ITP'13. Berlin, Heidelberg: Springer-Verlag. [https://doi.org/10.1007/978-3-642-39634-2\\_5](https://doi.org/10.1007/978-3-642-39634-2_5).
- Malecha, Gregory. 2014. "Extensible Proof Engineering in Intensional Type Theory." PhD thesis, Cambridge, Massachusetts: Harvard University, Graduate School of Arts & Sciences. <http://nrs.harvard.edu/urn-3:HUL.InstRepos:17467172>.

- McBride, Conor. 2005. “Epigram: Practical Programming with Dependent Types.” In *Advanced Functional Programming*, edited by Varmo Vene and Tarmo Uustalu, 3622:130–70. Berlin, Heidelberg: Springer Berlin Heidelberg. [https://doi.org/10.1007/11546382\\_3](https://doi.org/10.1007/11546382_3).
- Moura, Leonardo de, Soonho Kong, Jeremy Avigad, Floris van Doorn, and Jakob von Raumer. 2015. “The Lean Theorem Prover (System Description).” In *Automated Deduction - CADE-25*, edited by Amy P. Felty and Aart Middeldorp, 378–88. Lecture Notes in Computer Science. Cham: Springer International Publishing. [https://doi.org/10.1007/978-3-319-21401-6\\_26](https://doi.org/10.1007/978-3-319-21401-6_26).
- Moura, Leonardo de, and Sebastian Ullrich. 2021. “The Lean 4 Theorem Prover and Programming Language.” In *Automated Deduction - CADE 28*, edited by André Platzer and Geoff Sutcliffe, 625–35. Lecture Notes in Computer Science. Cham: Springer International Publishing. [https://doi.org/10.1007/978-3-030-79876-5\\_37](https://doi.org/10.1007/978-3-030-79876-5_37).
- Norell, Ulf. 2007. “Towards a Practical Programming Language Based on Dependent Type Theory.” PhD thesis, Göteborg, Sweden: Chalmers University of Technology and Göteborg University. <https://www.cse.chalmers.se/~ulfn/papers/thesis.pdf>.
- Odersky, Martin, Martin Sulzmann, and Martin Wehr. 1999. “Type Inference with Constrained Types.” *Theory and Practice of Object Systems* 5 (1): 35–55. [https://doi.org/10.1002/\(SICI\)1096-9942\(199901/03\)5:1%3C35::AID-TAPO4%3E3.0.CO;2-4](https://doi.org/10.1002/(SICI)1096-9942(199901/03)5:1%3C35::AID-TAPO4%3E3.0.CO;2-4).
- Orchard, Dominic, and Tom Schrijvers. 2010. “Haskell Type Constraints Unleashed.” In *Proceedings of the 10th International Conference on Functional and Logic Programming*, 56–71. FLOPS’10. Berlin, Heidelberg: Springer-Verlag. [https://doi.org/10.1007/978-3-642-12251-4\\_6](https://doi.org/10.1007/978-3-642-12251-4_6).
- Pédrot, Pierre-Marie. 2019. “Ltac2: Tactical Warfare.” In, 3. Cascais, Portugal. <https://popl19.sigplan.org/details/CoqPL-2019/8/Ltac2-Tactical-Warfare>.
- Peyton Jones, Simon. 2019. “Type Inference as Constraint Solving: How GHC’s Type Inference Engine Actually Works.” <https://www.microsoft.com/en-us/research/publication/type-inference-as-constraint-solving-how-ghcs-type-inference-engine-actually-works/>.
- Peyton Jones, Simon, Dimitrios Vytiniotis, Stephanie Weirich, and Mark Shields. 2007. “Practical Type Inference for Arbitrary-Rank Types.” *Journal of Functional Programming* 17 (1): 1–82. <https://doi.org/10.1017/S0956796806006034>.
- Pfenning, F. 1991. “Unification and Anti-Unification in the Calculus of Constructions.” In *[1991] Proceedings Sixth Annual IEEE Symposium on Logic in Computer Science*, 74–85. <https://doi.org/10.1109/LICS.1991.151632>.
- PTOOP. 2022. “Type Inference Is a Thought Trap. Write More Types and Fewer Programs.” Tweet. Twitter. February 7, 2022. <https://twitter.com/PTOOP/status/1490496253276340227>.
- Robert Estelle. 2019. “Heavy Coupling of Haskell Source Modules · Issue #3512 · Agda/Agda.” GitHub, Agda issues. January 20, 2019. <https://github.com/agda/agda/issues/3512>.
- Saibi, Amokrane. 1999. “Outils Génériques de Modélisation et de Démonstration pour la Formalisation des Mathématiques en Théorie des Types. Application à la Théorie des Catégories.” PhD thesis, Université Pierre et Marie Curie - Paris VI. <https://doi.org/document;https://web.archive.org/web/20230726150206/https://theses.hubscienc> 00523810.
- Spiwack, Arnaud. 2011. “Verified Computing in Homological Algebra.” PhD thesis, Ecole Polytechnique X. <https://pastel.archives-ouvertes.fr/pastel-00605836>.
- Swierstra, Wouter. 2008. “Data Types à La Carte.” *Journal of Functional Programming* 18 (4): 423–36. <https://doi.org/10.1017/S0956796808006758>.
- Tassi, Enrico, Claudio Sacerdoti Coen, Wilmer Ricciotti, and Andrea Asperti. 2012. “A Bi-Directional Refinement Algorithm for the Calculus of (Co)Inductive Constructions.” *Logical Methods in Computer Science* Volume 8, Issue 1 (March). [https://doi.org/10.2168/LMCS-8\(1:18\)2012](https://doi.org/10.2168/LMCS-8(1:18)2012).
- Tejiščák, Matúš. 2020. “A Dependently Typed Calculus with Pattern Matching and Erasure Inference.” *Proceedings of the ACM on Programming Languages* 4 (August): 91:1–29. <https://doi.org/10.1145/3408973>.
- The Agda Team. 2022. “Agda User Manual.” [https://agda.readthedocs.io/\\_/downloads/en/v2.6.2.2/pdf/](https://agda.readthedocs.io/_/downloads/en/v2.6.2.2/pdf/).
- The Coq Development Team. 2022. *The Coq Proof Assistant, Version 8.15.0*. Zenodo. <https://doi.org/10.5281/zenodo.5846982>.
- The Idris Team. 2021. “The Idris Tutorial.” [https://docs.idris-lang.org/\\_/downloads/en/v1.3.4/pdf/](https://docs.idris-lang.org/_/downloads/en/v1.3.4/pdf/).
- Ullrich, Sebastian, and Leonardo de Moura. 2020. “Beyond Notations: Hygienic Macro Expansion for Theorem Proving Languages.” In *Automated Reasoning*, edited by Nicolas Peltier and Viorica Sofronie-Stokkermans, 167–82. Lecture Notes in Computer Science. Cham: Springer International Publishing. [https://doi.org/10.1007/978-3-030-51054-1\\_10](https://doi.org/10.1007/978-3-030-51054-1_10).
- Víctor López Juan. 2020. *Tog+* (version 1.0). <https://framagit.org/vlopez/togt>.
- . 2021. “Practical Heterogeneous Unification for Dependent Type Checking.” Chalmers University of Technology. <https://research.chalmers.se/en/publication/527051>.
- Vytiniotis, Dimitrios, Simon Peyton Jones, Tom Schrijvers, and Martin Sulzmann. 2011. “OutsideIn(X) Modular Type Inference with Local Assumptions.” *Journal of Functional Programming* 21 (4-5): 333–412. <https://doi.org/10.1017/S0956796811000098>.
- Weirich, Stephanie. 2022. “Implementing Dependent Types in Pi-Forall.” OPLSS 2022: Lecture notes. <https://raw.githubusercontent.com/sweirich/pi-forall/2022/doc/oplss.pdf>.